

High-level concurrency concepts

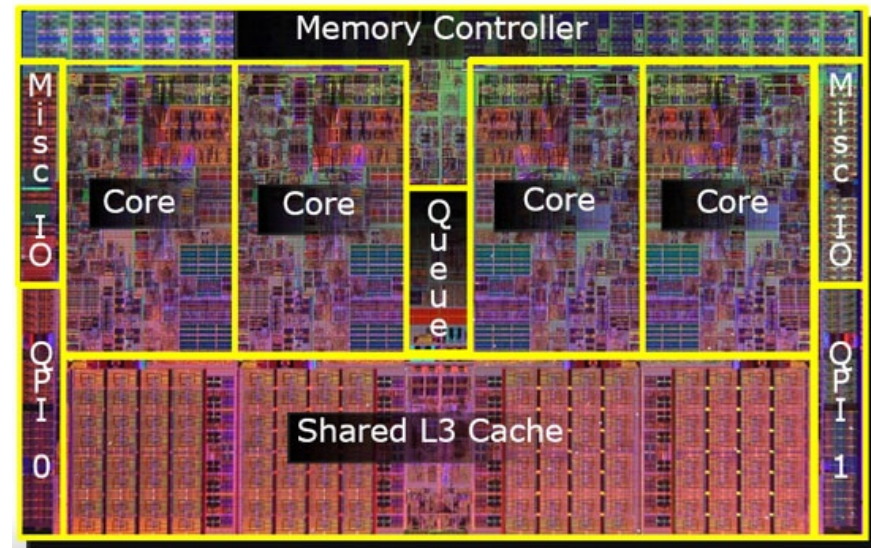


Václav Pech
NPRG014 2024/2025

<http://www.vaclavpech.eu>

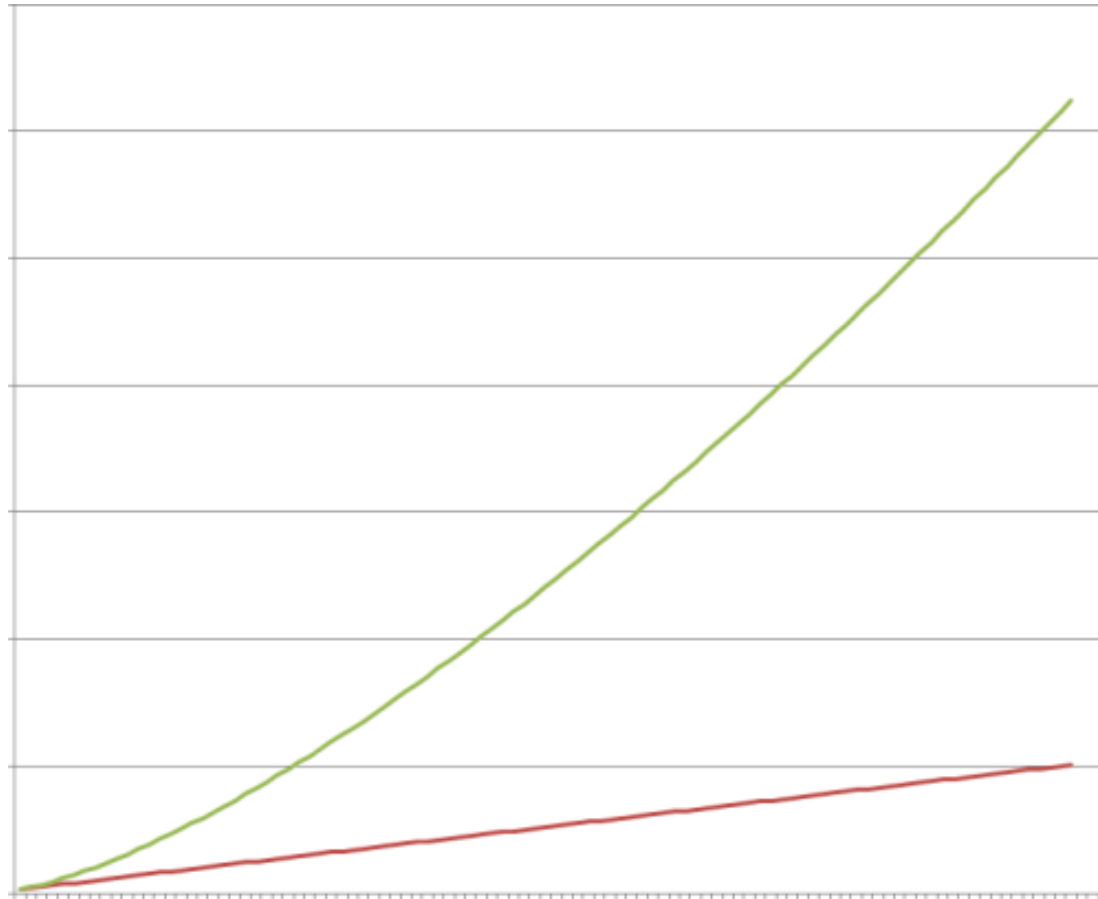
@vaclav_pech

Why concurrency?



We're all in the parallel computing business!

of cores



JVM low-level machinery

Thread, Runnable, Thread Pools, Virtual
Threads

– to run things in parallel

JVM low-level machinery

Thread, Runnable, Thread Pools, Virtual Threads

– to run things in parallel

Synchronized blocks

Volatile

Locks

Atomic

– to run things sequentially again

Threads

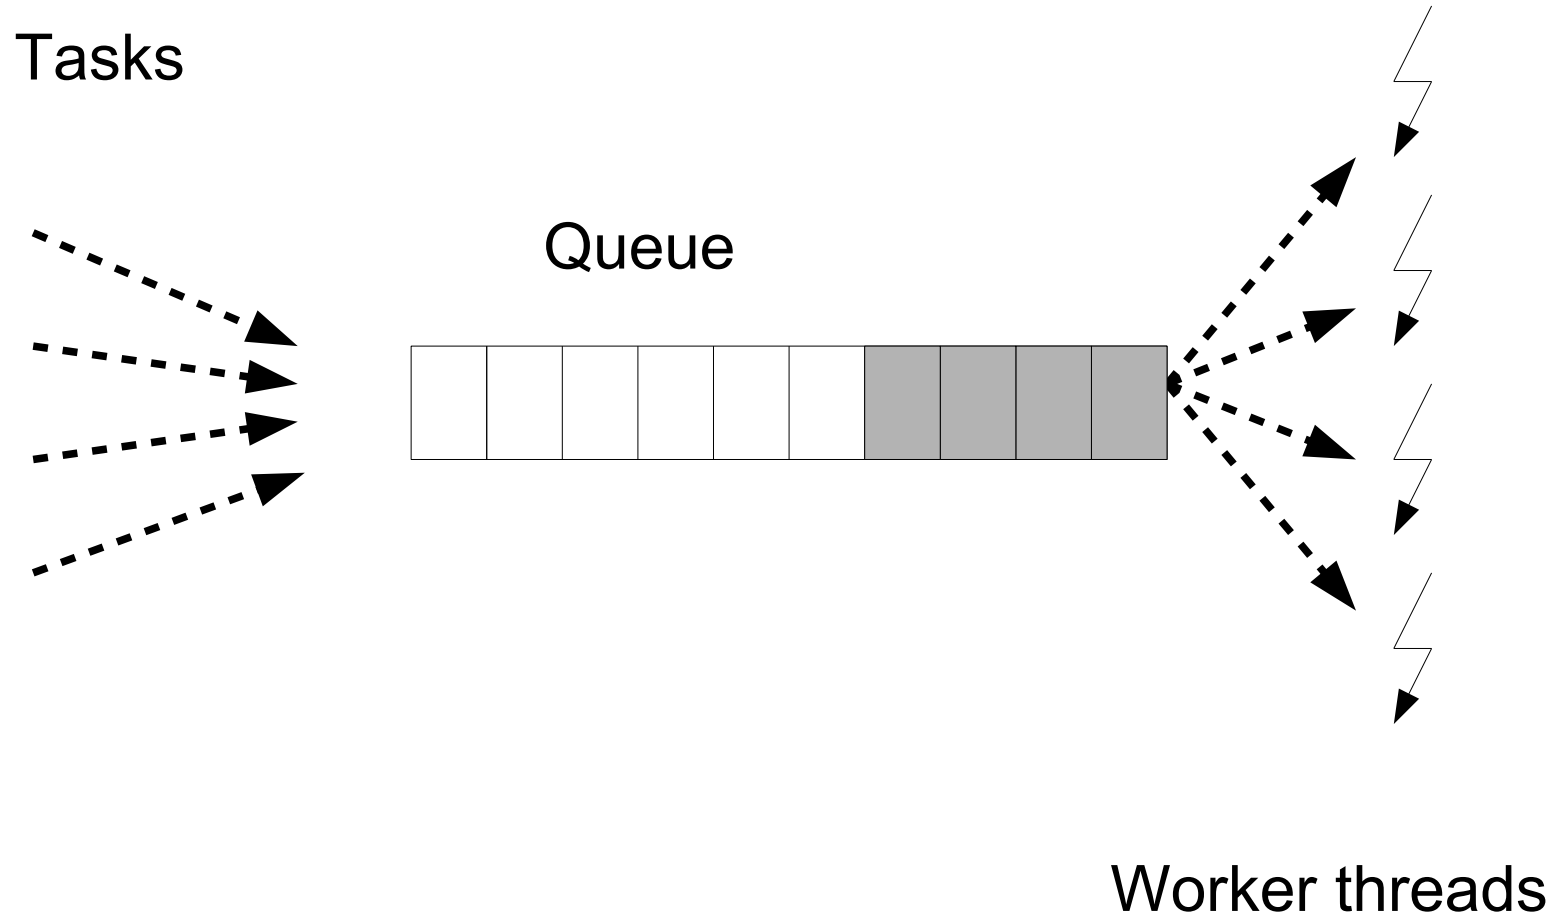
```
def myRunnable = new Runnable {  
    public void run() {  
        ...  
    }  
}  
  
new Thread(myRunnable).start()
```

Thread Pool

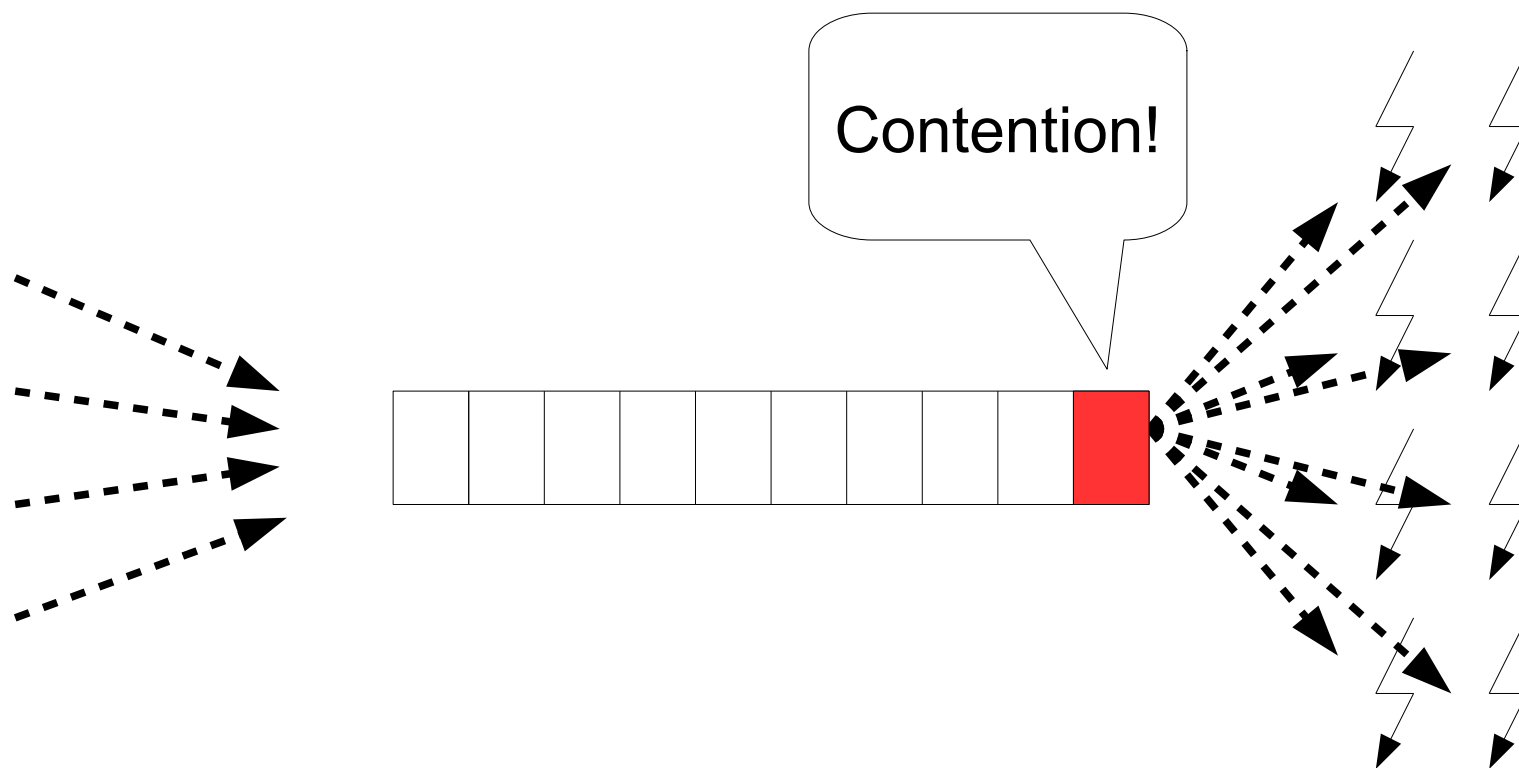
```
def pool = Executors.newFixedThreadPool(6)
def myCallable = new Callable<T> {
    public T call() {...}
}
```

```
Future<T> f = pool.submit(myCallable)
println f.get()
```


Thread Pool



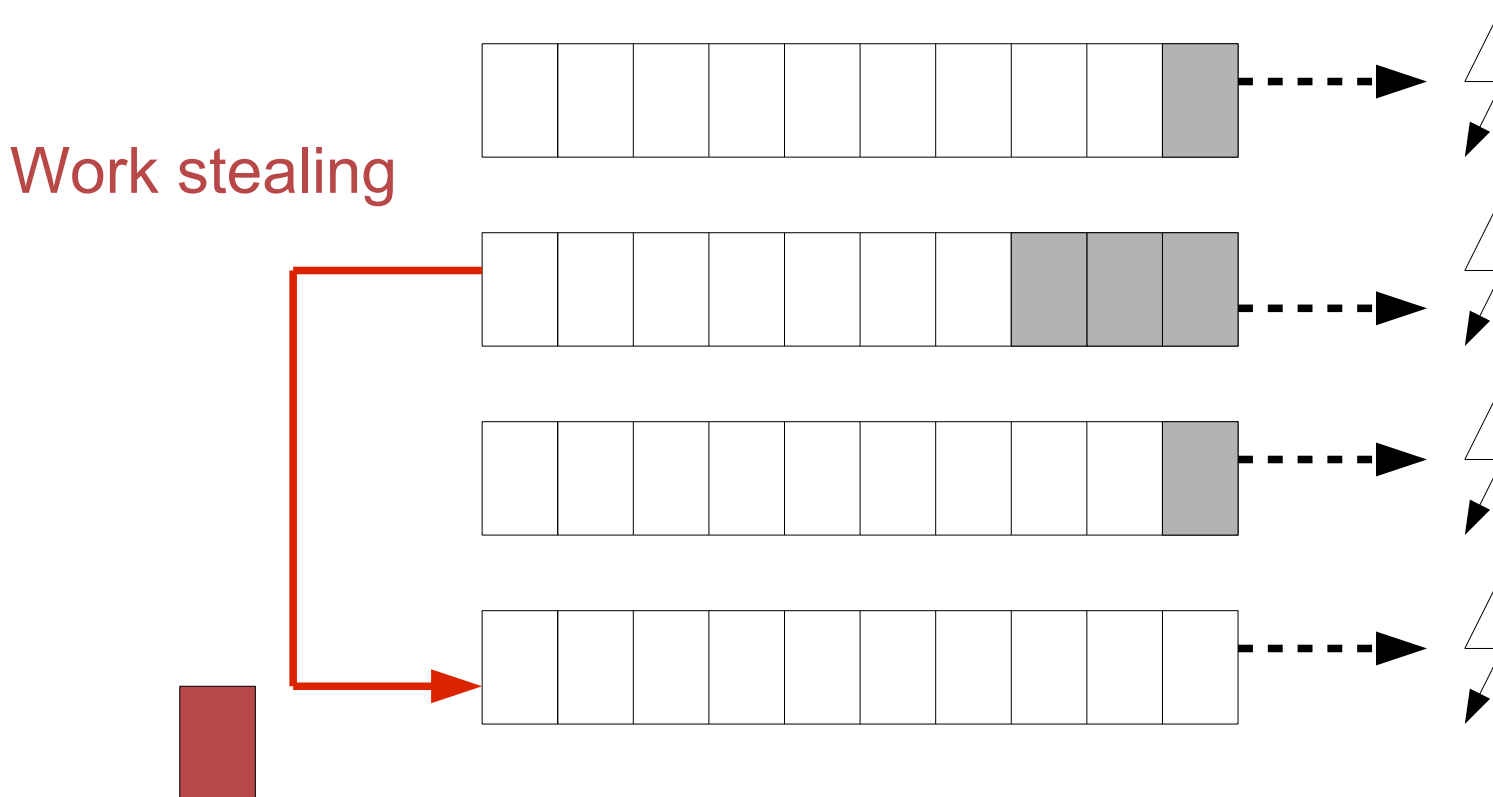
Thread Pool



Fork/Join Thread Pool



Fork/Join Thread Pool



Virtual Threads (Java 19)

```
def t = Thread.ofVirtual.start(myRunnable)  
println f.join()
```

Virtual Threads (Java 21)

```
def t = Thread.ofVirtual.start(myRunnable)  
println f.join()
```

```
def pool = Executors  
.newVirtualThreadPerTaskExecutor()  
Future<T> f = pool.submit(myCallable)  
println f.get()
```

Scopes (Java 21)

```
Response handle() throws ExecutionException, InterruptedException {  
    try (var scope = new StructuredTaskScope.ShutdownOnFailure()) {  
        Supplier<String> user = scope.fork(() -> findUser());  
        Supplier<Integer> order = scope.fork(() -> fetchOrder());  
  
        scope.join()           // Join both subtasks  
            .throwIfFailed(); // ... and propagate errors  
  
        // Here, both subtasks have succeeded, so compose their results  
        return new Response(user.get(), order.get());  
    }  
    //...  
}
```

Scope

Provides virtual threads with:

- a thread pool
- a context
- *life-cycle management*
- *control for sub-groups*
- *error handling for tasks and sub-groups*

Thread pools the Groovy way

```
task {  
    calculation.process()  
}
```



Scopes the Groovy way

```
def group = new NonDaemonPGroup(10)
```

```
group.task {  
    calculation.process()  
}
```



Task alternatives

```
group.task {->...}
```

```
group.task new Runnable() {...}
```

```
group.task new Callable<V>() {...}
```



Not implemented yet

`group.cancelAll()`

`group.onFailureRestart()`

`group.onFailureCancelAll()`

`def subgroup = group.subgroup()`



Dealing with threads sucks!

```
public class Counter {  
    private static long count = 0;  
  
    public Counter() {  
  
        count++;  
  
    }  
}
```

Dealing with threads sucks!

```
public class Counter {  
    private volatile static long count = 0;  
  
    public Counter() {  
  
        count++;  
  
    }  
}
```

Dealing with threads sucks!

```
public class Counter {  
    private volatile static long count = 0;  
  
    public Counter() {  
  
        count = count + 1;  
  
    }  
}
```

Dealing with threads sucks!

```
public class Counter {  
    private static long count = 0;  
  
    public Counter() {  
        synchronized (this) {  
            count++;  
        }  
    }  
}
```

Dealing with threads sucks!

```
public class Counter {  
    private static long count = 0;  
  
    public Counter() {  
        synchronized (this.getClass()) {  
            count++;  
        }  
    }  
}
```


Dealing with threads sucks!

```
public class Counter {  
    private static Long count = 0;  
  
    public Counter() {  
        synchronized (count) {  
            count++;  
        }  
    }  
}
```



Dealing with threads sucks!

```
public class Counter {  
    private static Long count = 0;  
  
    public Counter() {  
        synchronized (count) {  
            count = new Long(count.longValue() + 1);  
        }  
    }  
}
```

Dealing with threads sucks!

```
public class ClickCounter implements ActionListener {  
    private int numberOfClicks;  
    public ClickCounter(JButton button, int initValue) {  
        button.addActionListener(this);  
        numberOfClicks = initValue;  
    }  
    public void actionPerformed(final ActionEvent e) {  
        numberOfClicks++;  
    }  
}
```

Dealing with threads sucks!

```
public class ClickCounter implements ActionListener {  
    private int numberOfClicks;  
    public ClickCounter(JButton button, int initValue) {  
        button.addActionListener(this);  
        numberOfClicks = initValue;   
    }  
    public void actionPerformed(final ActionEvent e) {  
        numberOfClicks++;   
    }  
}
```

Low-level concurrency

- Monitors
- Locks – R/W, ReentrantLock
- Semaphors
- CountdownLatches
- Barriers

Stone age of parallel SW


Dead-locks

Live-locks

Race conditions

Starvation

Shared Mutable State



Why high-level concurrency?

Multithreaded programs today work mostly by accident!



The root problem

Shared mutable state

Approach 1

Shared **mutable** state

Use read-only data

Approach 2

Shared mutable state

Use only local thread-contained data

High-level abstractions

- STM
- Hierarchical decomposition
- Geometrical decomposition
- Streams
- Agents, Actors, Active objects
- Communicating Sequential Processes

High-level abstractions

- STM
- Hierarchical decomposition
- Geometrical decomposition
- Streams
- Agents, Actors, Active objects
- Communicating Sequential Processes


Control State Sharing

```
List registrations = []  
submissions.each {form →  
  group.task {  
    if (form.process().valid) {  
      registrations << form  
    }  
  }  
}  
}
```

Control State Sharing

Needs protection

```
List registrations = []  
submissions.each {form →  
  group.task {  
    if (form.process().valid) {  
      registrations << form  
    }  
  }  
}
```



STM (Akka - Scala)

```
atomic {  
  .. // do something within a transaction  
}
```

```
atomic(maxNrOfRetries) { .. }  
atomicReadOnly { .. }
```

```
atomically {  
  .. // try to do something  
} orElse {  
  .. // if tx clash; try do do something else  
}
```

Persistent Data Structures

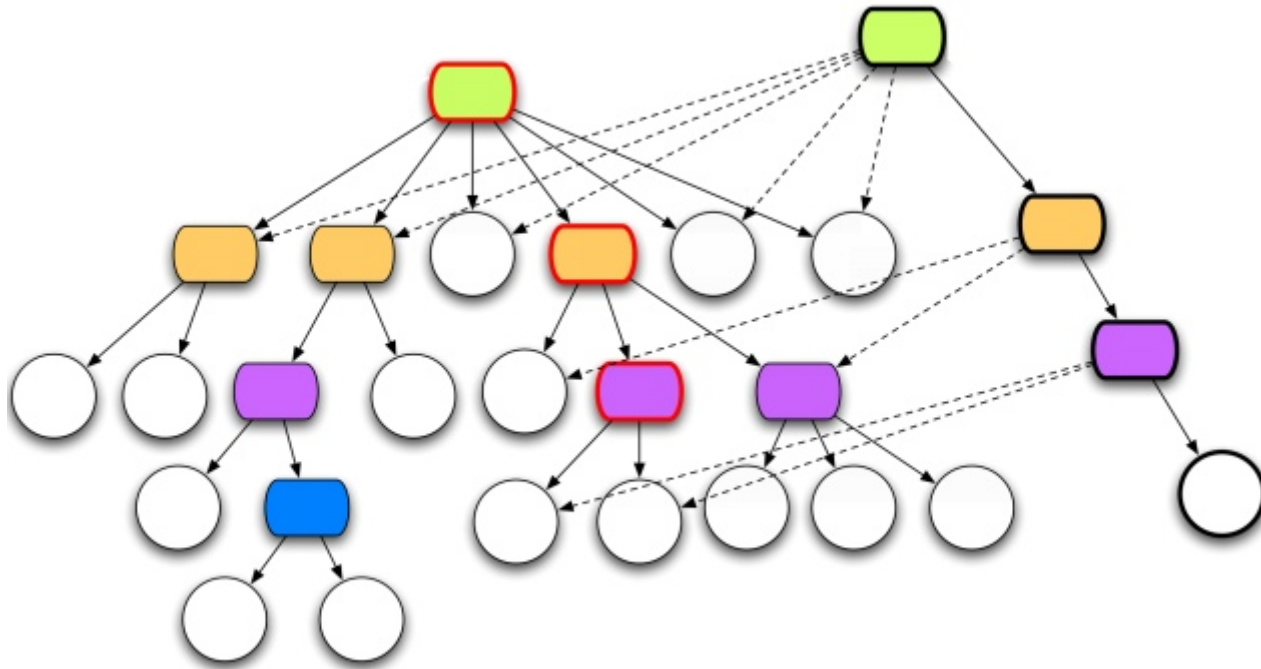


Illustration taken from Rich Hickey's presentation. Copyright Rich Hickey 2009

High-level abstractions

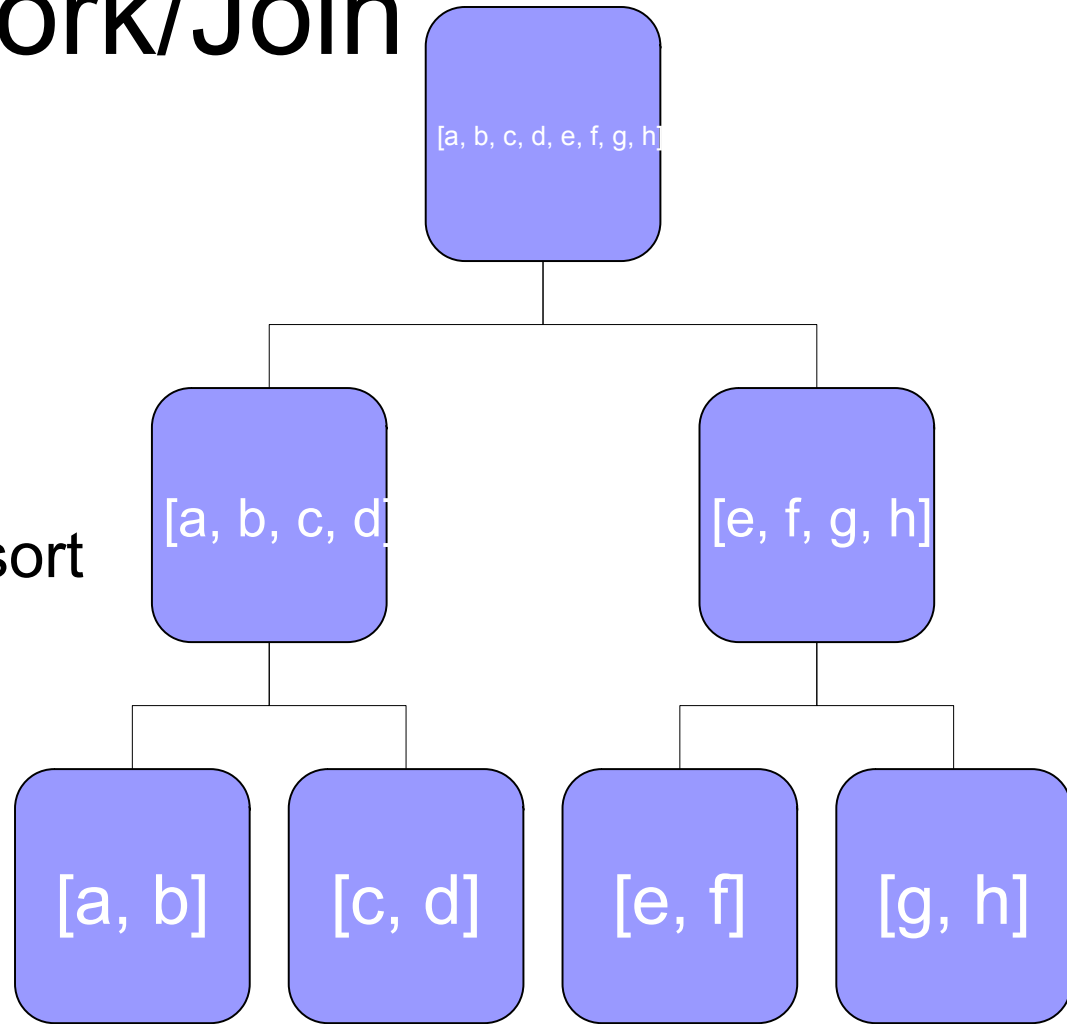
- STM
- Hierarchical decomposition
- Geometrical decomposition
- Streams
- Agents, Actors, Active objects
- Communicating Sequential Processes

Hierarchical decomposition

[64, 63, 62, 61, 60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50, 49, 48, 47, 46, 45, 44, 43, 42, 41, 40, 39, 38, 37, 36, 35, 34, 33, 32, 31, 30, 29, 28, 27, 26, 25, 24, 23, 22, 21, 20, 19, 18, 17, 16, 15, 14, 13, 12, 11, 10, 9, 8, 7, 6, 5, 4, 3, 2, 1]									
[64, 63, 62, 61, 60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50, 49, 48, 47, 46, 45, 44, 43, 42, 41, 40, 39, 38, 37, 36, 35, 34, 33]								[32, 31, 30, 29, 28, 27, 26, 25, 24, 23, 22, 21, 20, 19, 18,	
[64, 63, 62, 61, 60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50, 49]				[48, 47, 46, 45, 44, 43, 42, 41, 40, 39, 38, 37, 36, 35, 34, 33]				[32, 31, 30, 29, 28, 27, 26, 25, 24, 23, 22, 21, 20, 19, 18,	
[64, 63, 62, 61, 60, 59, 58, 57]		[56, 55, 54, 53, 52, 51, 50, 49]		[48, 47, 46, 45, 44, 43, 42, 41]		[40, 39, 38, 37, 36, 35, 34, 33]			
	[56, 55, 54, 53]		[49, 50, 51, 52]			[40, 39, 38, 37]		[33, 34, 35, 36]	
	[56, 55]		[51, 52] [49, 50]			[40, 39]		[35, 36] [33, 34]	

Fork/Join

- Solve hierarchical problems
 - Divide and conquer
 - Merge sort, Quick sort
 - Tree traversal
 - File scan / search
 - ...



Fork/Join (GPars)

```
{currentDir ->
  long count = 0;
  currentDir.eachFile {
    if (it.isDirectory()) {
      forkOffChild it
    } else {
      count++
    }
  }
  return count + childrenResults.sum(0)
}
```

Waits for children
without blocking the
thread!



High-level abstractions

- STM
- Hierarchical decomposition
- Geometrical decomposition
- Streams
- Agents, Actors, Active objects
- Communicating Sequential Processes

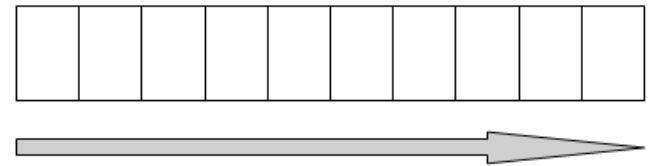
Geometric decomposition

images.`eachParallel` {it.process()}

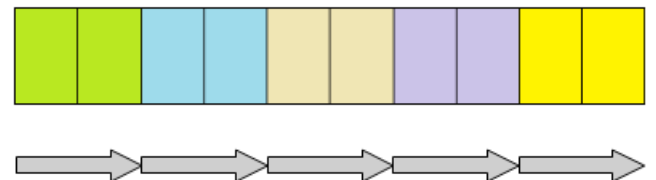
documents.`sumParallel`()

candidates.`maxParallel` {it.salary}.marry()

1 thread



5 threads



Geometric decomposition

```
registrations = submissions
```

```
  .collectParallel { form -> form.process() }
```

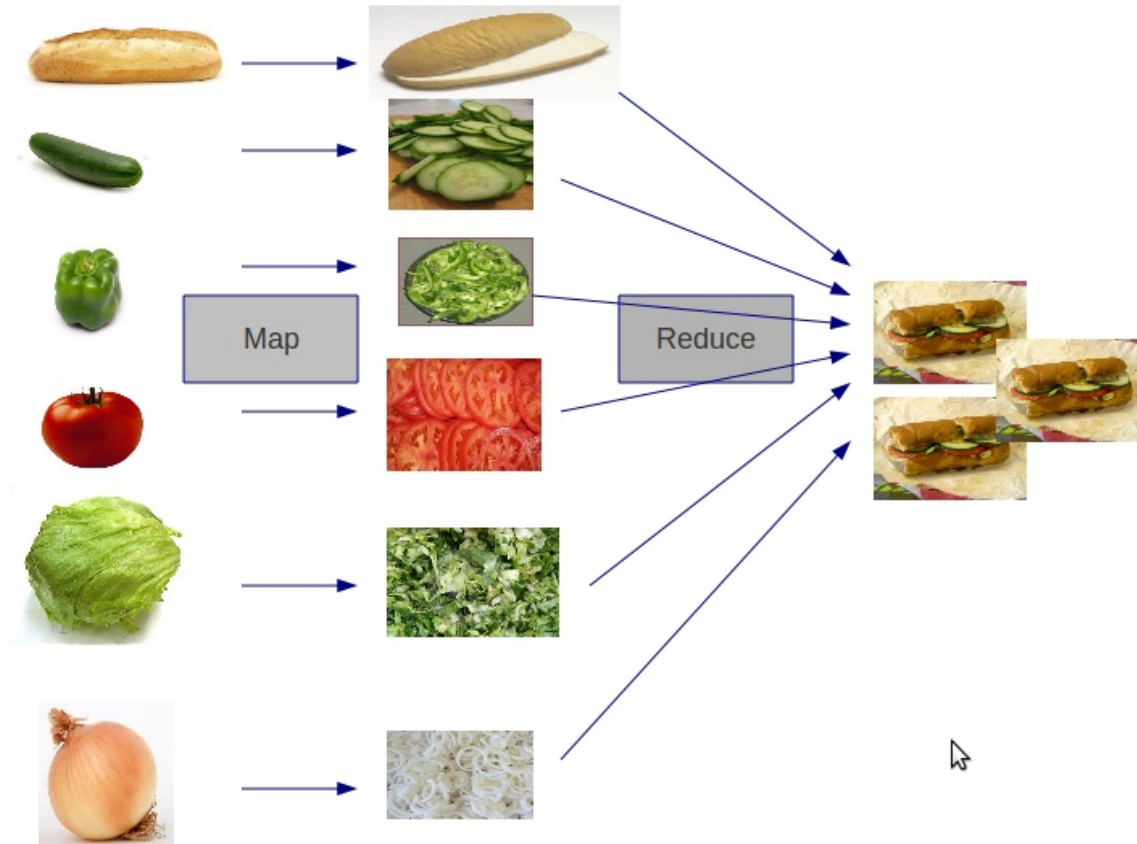
```
  .findAllParallel { it.valid }
```

```
registrations = submissions.parallel
```

```
  .map { form -> form.process() }
```

```
  .filter { it.valid }.collection
```

Map - reduce




Frequent confusion

[Questions](#)[Tags](#)[Users](#)[Badges](#)[Unanswered](#)

parallel quick sort outdone by single threaded quicksort

- ▲ I've been reading , here is the example in the book using futures to implement parallel quick sort.
 - 0 But I found this function is more than twice slower than the single threaded quick sort function without using any asynchronous facilities in c++ standard library. Tested with g++ 4.8 and visual c++ 2012.
 - ▼
 - ☆ I used 10M random integers to test, and in visual c++ 2012, this function spawned 6 threads in total to perform the operation in my quad core PC.
- I am really confused about the performance. Any body can tell me why?

GPU



Improper use 1

```
def accumulator = 0
```

```
myCollection.eachParallel {  
    accumulator += calculate(it)  
}
```

Do not accumulate, map-reduce!

```
def accumulator = myCollection.parallel  
    .map {calculate(it)}  
    .reduce {a, b → a + b}
```

Improper use 2

```
new File("/file.txt").withReader{reader ->
  reader.eachParallel {
    def r1 = step1(r)
    def r2 = step2(r1)
    def r3 = step3(r2)
  }
}
```

Unroll iterations into streams

```
def pipeline = data | step1 | step2 | step3
```

```
new File("/file.txt").withReader{reader ->
  reader.each {
    data << it
  }
}
```


Unroll iterations

```
[project1, project2, project3, project4].eachParallel {  
  checkout it  
  compile it  
  package it  
  deploy it  
}
```

vs.

```
def results = projects | checkout | compile | package | deploy
```

Unroll iteration



High-level abstractions

- STM
- Hierarchical decomposition
- Geometrical decomposition
- **Streams**
- Agents, Actors, Active objects
- Communicating Sequential Processes


Pipeline DSL

```
def toProcess = new DataflowQueue()
```

```
def validated = new DataflowQueue()
```

```
toProcess | {form -> process(form)} |  
    {processedForm -> validate(processedForm)} | validated
```


```
submissions.each {toProcess << it}
```



Generalization

An operation can take multiple inputs and outputs

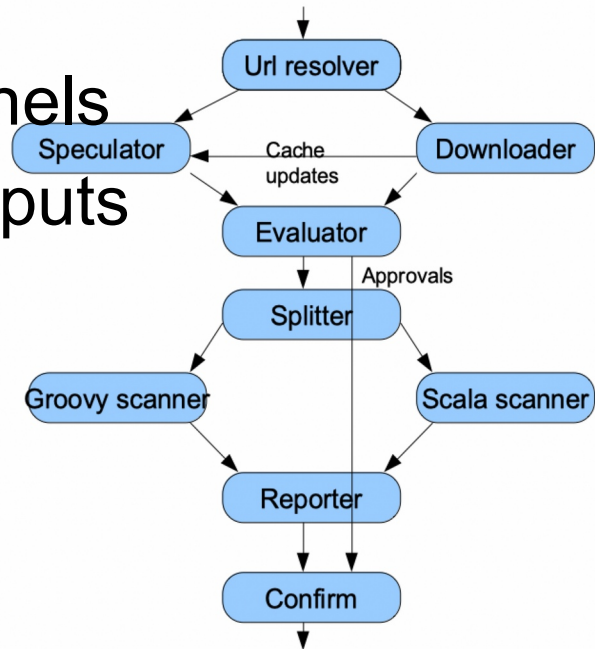
Operations addressed indirectly using channels



Dataflow Operators

Like functions

- wired together with channels
- possibly with multiple outputs
- may be recursive



Dataflow Operators

```
operator(inputs: [headers, bodies, footers],  
         outputs: [articles, summaries])
```

```
{header, body, footer ->
```


```
  def article = buildArticle(header, body, footer)
```



```
  bindOutput(0, article)
```


```
  bindOutput(1, buildSummary(article))
```

```
}
```

Factorial








High-level abstractions


- STM
- Hierarchical decomposition
- Geometrical decomposition
- Streams
- **Agents, Actors, Active objects**
- Communicating Sequential Processes

Agent

Lock **Shared Mutable State** in a **Safe**



Agent inside



Sharing through agents

```
Agent registrations = new Agent( [] )
```

```
submissions.each {form →
```

```
  task {
```

```
    if (form.process().valid) {
```

```
      registrations.send {it << form}
```

```
    }
```

```
  }
```

```
}
```

Actors

Agents send commands

Actors send data

Both use direct addressing (no channels)

Actors

Isolated state


Active - shared threads

Async Communication

- Direct addressing
- Immutable messages

Activities:

- Create a new actor
- Send a message
- Receive a message



Actor definition

```
class MyActor extends DynamicDispatchActor {  
    private int counter = 0  
  
    public void onMessage(String msg) {  
        this.counter += msg.size()  
    }  
    public void onMessage(Integer number) {  
        this.counter += number  
    }  
    public void onMessage(Money cash) {  
        this.counter += cash.amount  
        reply 'Thank you'  
    }  
}
```

Sending messages


```
buddy.send 10.eur
```

```
buddy << new Book(title:'Groovy Recipes',  
                    author:'Scott Davis')
```

```
def canChat = buddy.sendAndWait 'Got time?'
```

```
buddy.sendAndContinue 'Need money!', {cash->  
    pocket.add cash  
}
```

Actors use



Enricher
Router
Translator
Endpoint
Splitter
Aggregator
Filter
Re-sequencer
Checker

Router

Endpoint

Splitter

Aggregator

Filter

Re-sequencer

Checker

Active objects

@ActiveObject

```
class MyCounter {  
    private int counter = 0
```

@ActiveMethod

```
    def incrementBy(int value) {  
        println "Received an integer: $value"  
        this.counter += value  
    }  
}
```

High-level abstractions

- STM
- Hierarchical decomposition
- Geometrical decomposition
- Streams
- Agents, Actors, Active objects
- Communicating Sequential Processes

CSP

Communicating Sequential Processes

Focus on composable processes more than on data

Abstractions

Tasks (aka coroutines, goroutines)

- Implemented as (pooled) threads or virtual, lightweight threads

Data exchange between tasks

- Tasks' *promised* return value
- Dataflow variables
- Dataflow channels

Tasks


```
requests.each {r ->  
  task {  
    r.handle()  
    println "Done processing " + r.id  
  }  
}
```

Dataflow Variables


No race-conditions

No live-locks

Deterministic deadlocks



Dataflow Variables / Promises



DataflowVariables/Promises

```
def x = new DataflowVariable()
```

```
task { println x.get() }
```

```
task {  
    storeInDB(x.get())  
}
```

```
task { x << 10 }
```

```
x.join()
```

Promises to exchange data

```
def x,y,z = new DataflowVariable()
```

```
task { z << x.get() + y.get() }
```

```
task { x << 10 }
```

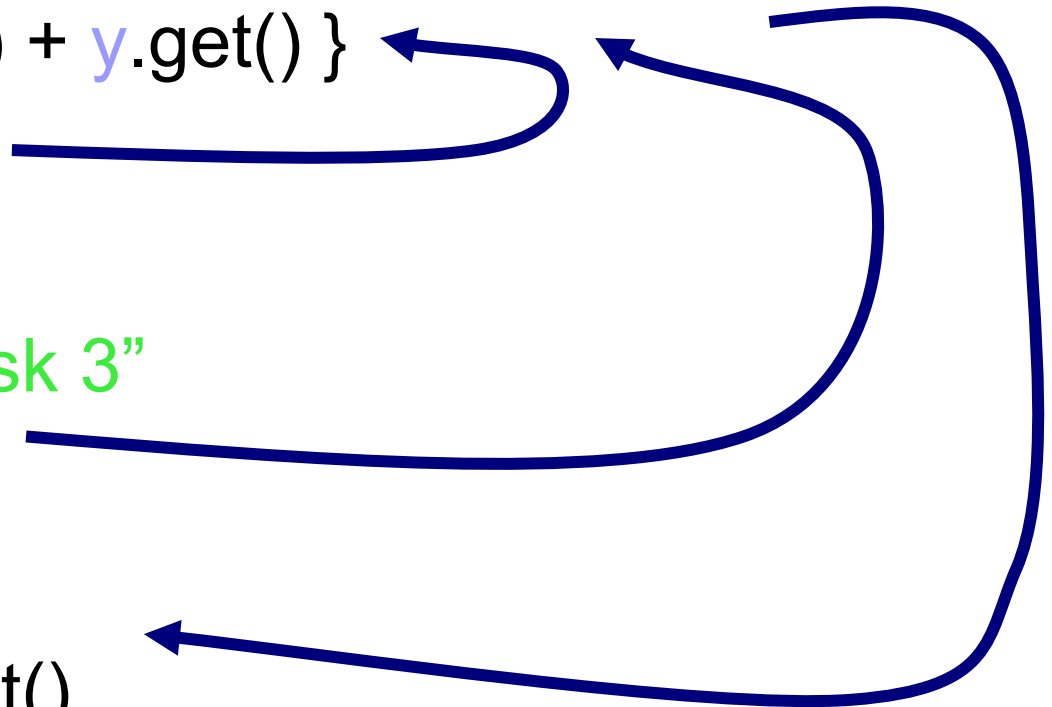
```
task {
```

```
    println "I am task 3"
```


```
    y << 5
```

```
}
```

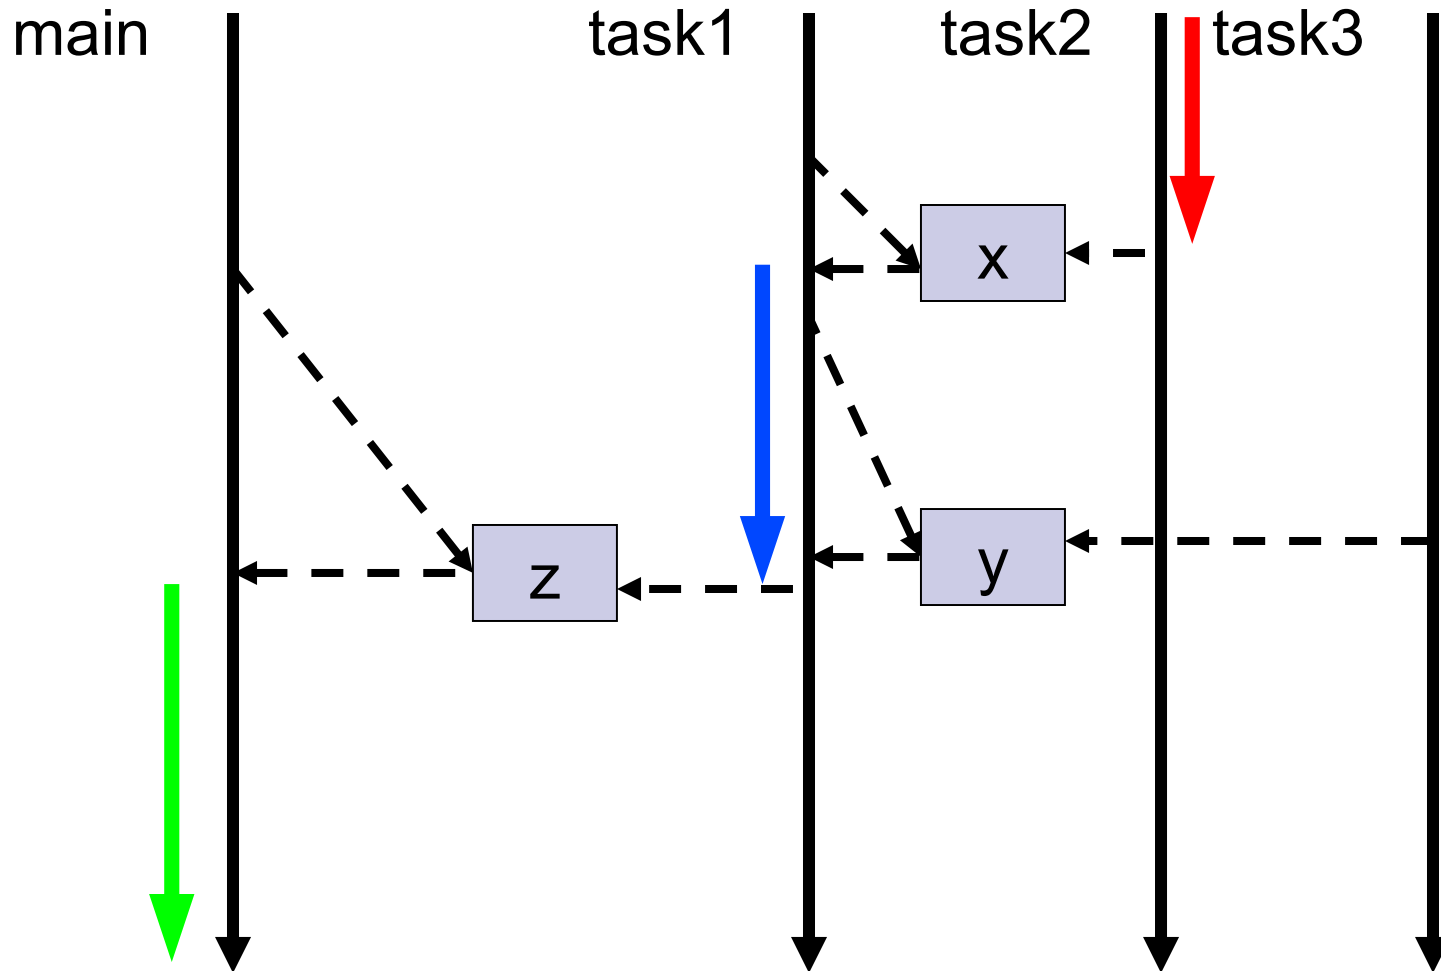
```
assert 15 == z.get()
```



Dataflow Variables / Promises



Dataflow Variables / Promises



Promises as return values

```
Promise x = task { return 10 }
```

```
Promise y = task {  
    println "I am task 3"  
    return 5  
}
```

```
Promise z = task {x.get() + y.get() }  
assert 15 == z.get()
```


Promises as return values

```
task { println "Hello" }
```

```
Promise t1 = task {  
    def data = loadDataFromDB("Joe")  
    data.optimize()  
    return data  
}
```

```
Promise t2 = task { printData(t1.get()) }  
t2.join()
```

Kotlin coroutines

+

```
val time = measureTimeMillis {  
    val one = async { doSomethingUsefulOne() }  
    val two = async { doSomethingUsefulTwo() }  
    println("The answer is ${one.await() + two.await()}")  
}  
println("Completed in $time ms")
```



Reactive approach

- Needed when tasks are implemented on top of system threads (*unlike virtual/lightweight threads*)
- Blocking in *Promise.get()* is too expensive

Reactive allows for callbacks to be registered on Promises

Callbacks

```
def x = new DataflowVariable()

x.then { println it }

task {
    x.then { storeInDB(it) }
        .then { dbStatus → println dbStatus }
}

task { x << 10 }

x.join()
```

Chaining promises

```
def h1 = download('url') then {text → text.trim()} then hash
```

Chaining promises

```
def h1 = download('url') then {text → text.trim()} then hash
```

```
def h1 = download('url') | {text → text.trim()} | hash
```

Error handling

```
url.then(download)  
    .then(calculateHash)  
    .then(formatResult)  
    .then(printResult, printError)  
    .then(sendNotificationEmail);
```

CompletableFuture (Java)

```
final CompletableFuture<Integer> f = CompletableFuture.supplyAsync(() -> 10);
```

```
final CompletableFuture<Integer> f2 = f.thenApplyAsync(x -> 2 * x)  
    .handleAsync((x, e) -> 3 * x);
```

```
final CompletableFuture<Integer> f3 = f.thenComposeAsync(x -> CompletableFuture.supplyAsync(() -> 2 * x))  
    .handleAsync((x, e) -> 3 * x);
```

```
final CompletableFuture<Integer> f4 = f2.thenCombineAsync(f3, (x, y) -> x + y);
```


Glue tasks together

```
Promise c1 = task { compile(module1) }
```

```
Promise c2 = task { compile(module2) }
```

Glue tasks together

```
Promise c1 = task { compile(module1) }
```

```
Promise c2 = task { compile(module2) }
```

```
Promise j1 = c1.then {jar it}
```

```
Promise j2 = c2.then {jar it}
```

Glue tasks together

```
Promise c1 = task { compile(module1) }
```

```
Promise c2 = task { compile(module2) }
```

```
Promise j1 = c1.then {jar it}
```

```
Promise j2 = c2.then {jar it}
```

```
whenAllBound(j1, j2) {m1, m2 → deploy(m1, m2)}
```

```
j1.then {pushToRepo it}
```

Glue tasks together

```
Promise c1 = task { compile(module1) }
```

```
Promise c2 = task { compile(module2) }
```

```
Promise j1 = c1.then {jar it}
```

```
Promise j2 = c2.then {jar it}
```

```
whenAllBound(j1, j2) {m1, m2 → deploy(m1, m2)}
```

```
j1.then {pushToRepo it}
```

```
iWillSendEmailWhenJarred(j1)
```

Glue tasks without callbacks

```
Promise c1 = task { compile(module1) }
```

```
Promise c2 = task { compile(module2) }
```

```
Promise j1 = task { jar c1.get() }
```

```
Promise j2 = task { jar c2.get() }
```

```
task { deploy(j1.get(), j2.get()) }
```

```
task {pushToRepo j1.get() }
```

```
iWillSendEmailWhenJarred(j1.get())
```

Dataflow variables wrap-up

- *Promise* is the reading side
- Blocking read – *val*, *get()*, *await()*
- Callbacks – *then*, *|*, *whenAllBound*
- Callback handlers can be chained
- Promises are typically returned from async calls – *task*, *async*

Abstractions


Tasks (aka coroutines, goroutines)

- Implemented as (pooled) threads or virtual, lightweight threads


Data exchange between tasks

- Tasks' return value
- Dataflow variables
- **Dataflow channels**


Dataflow Variables / Promises



Dataflow Channels



Synchronous Channels



Tasks as processes (CSP)

```
group.task {  
  doStuff()  
  logChannel << 'initialized'  
  def result = doWork(workQueue.get())  
  if (result.isError) errors << result  
  else results << result  
  logChannel << 'finished'  
}
```

Channel Selection

```
Select alt = group.select(validForms, invalidForms)
```

```
SelectResult selectResult = alt.select() //alt.prioritySelect()
```

```
switch (selectResult.index) {  
    case 0: registrations << selectResult.value; break  
    case 1: ...  
}
```

High-level abstractions

- STM
- Hierarchical decomposition
- Geometrical decomposition
- Streams
- Agents, Actors, Active objects
- Communicating Sequential Processes

Summary

Parallelism is not hard, multi-threading is

References

<http://groovy-lang.org>

<http://grails.org>

<http://groovyconsole.appspot.com/>

<http://www.manning.com/koenig2/>

<https://docs.oracle.com/en/java/javase/20/core/virtual-threads.html>

<https://akka.io/>